

A Step-by-Step Guide to 2D Motion Profiling Using Bézier Curves

Mechanics Division, Autonomous Systems Research

May 18, 2025

Abstract

This paper walks through the construction and explanation of a real-time motion profiling system for differential drive robots. It introduces foundational concepts in calculus, curve modeling, numerical integration, and motion planning. All concepts are broken down for high-school level understanding, with annotated code and diagrams to guide the reader through the complete system.

Contents

1	What is Motion Profiling?	3
2	Our Goal: Move Along a Curve	3
3	Curve Representation Options	3
4	What are Keyframes?	4
5	Quick Primer: Derivatives and Integrals	4
6	Arc Length via Gaussian Quadrature	4
7	Finding the Next Parameter: Newton–Raphson	6
8	Acceleration and Deceleration Distances	6
9	Curvature and Turning Speed	7
10	Keyframe Velocity Interpolation	7
11	Velocity Planning Algorithm	8
12	Handling Multiple Path Segments	8
13	Simplified C++ Example	8
14	Why Acceleration/Deceleration Distance is an Approximation	10
15	Ideas for Future Improvement	11
16	Further Learning Resources	11

17 Conclusion**11**

1 What is Motion Profiling?

Motion profiling answers the question: how should a robot move from point A to point B smoothly and safely?

- How fast should it go?
- When should it turn?
- When should it start slowing down?

The result is a path plus a time-parameterized velocity and angular velocity profile that respects the robot's physical limits.

2 Our Goal: Move Along a Curve

We want to follow a curved path while:

- Respecting maximum speed and acceleration
- Slowing down for turns (based on curvature)
- Hitting specified speeds at key positions (keyframes)

To do that, we describe the path mathematically and then compute how the robot should move along it in real time.

3 Curve Representation Options

There are multiple ways to describe a smooth curve:

- **Quadratic Bézier Curve:** Uses 3 control points (a start point, end point, and one middle control point) to form a smooth curve. It is simple but less flexible in shaping the path.
- **Cubic Bézier Curve:** Uses 4 control points (start, end, and two control points). This gives more control over the curve's shape. We use cubic Béziers here for a good balance of simplicity and flexibility.
- **Cubic Hermite Splines:** Defined by endpoints and specified tangents (slopes) at those endpoints. This means you explicitly set the direction of travel at each end, resulting in a smooth curve that honors those directions.
- **Catmull–Rom Splines:** A type of spline that passes through a series of given waypoints. The curve is calculated to smoothly interpolate through all the points, making it easy to create paths through specific locations.

We chose cubic Bezier for its balance of simplicity and control.

4 What are Keyframes?

Keyframes are specific points along the path where we fix the robot's speed. For example:

- Start at 0 m/s
- Reach 1 m/s halfway
- End at 0 m/s

The planner converts those into arc-length positions and then blends between them. In practice, path planning software (such as *PATH.JERRYIO*) provides a visual interface to define keyframes. For example, *PATH.JERRYIO* includes a "speed graph" where you can add **speed keyframes** at various positions along the path: `contentReference[oaicite:1]index=1`. Each keyframe corresponds to a specific distance along the path (the x-axis on the graph) and a desired speed (the y-axis). The planner then constructs the speed profile by connecting these points smoothly, ensuring the robot slows down or speeds up at the designated locations.

5 Quick Primer: Derivatives and Integrals

Derivative = Rate of Change

If $s(t)$ is position over time, then

$$v(t) = \frac{ds}{dt},$$

the instantaneous speed.

Integral = Area Under Curve

If $v(t)$ is speed, then

$$s(t) = \int_0^t v(\tau) d\tau,$$

the total distance traveled.

6 Arc Length via Gaussian Quadrature

For a parametric curve $\mathbf{r}(t) = (x(t), y(t))$, the arc length from the start up to parameter t is

$$s(t) = \int_0^t \|\mathbf{r}'(\tau)\| d\tau,$$

the integral of speed $\|\mathbf{r}'(\tau)\|$ along the curve. We approximate this integral with Gaussian quadrature:

$$\int_a^b f(x) dx \approx \frac{b-a}{2} \sum_{i=1}^n w_i f\left(\frac{b-a}{2}x_i + \frac{a+b}{2}\right).$$

Numerical integration rationale. For most curves, there isn't a simple formula for $s(t)$, so we need to integrate numerically. Gaussian quadrature is one effective method for numerical integration. It chooses special sample points (nodes x_i) and weights w_i to maximize accuracy. An n -point Gauss–Legendre rule is exact for all polynomials up to degree $2n - 1$, meaning it integrates any such polynomial perfectly on $[-1, 1]$ with only n evaluations of f .

To illustrate, the 2-point Gaussian rule on $[-1, 1]$ picks

$$x_1 = -\frac{1}{\sqrt{3}}, \quad x_2 = \frac{1}{\sqrt{3}}, \quad w_1 = w_2 = 1,$$

so that

$$\int_{-1}^1 f(x) dx \approx f\left(-\frac{1}{\sqrt{3}}\right) + f\left(\frac{1}{\sqrt{3}}\right),$$

which integrates any cubic polynomial exactly.

Table of Nodes and Weights. The following table lists the Gauss–Legendre nodes x_i and weights w_i on the standard interval $[-1, 1]$ for $n = 2, 3, 4$. These are the constants you plug into the formula above.

n	x_i (nodes)	w_i (weights)
2	$-\frac{1}{\sqrt{3}} \approx -0.577350$	1
	$+\frac{1}{\sqrt{3}} \approx 0.577350$	1
3	$-\sqrt{\frac{3}{5}} \approx -0.774597$	$\frac{5}{9} \approx 0.555556$
	0	$\frac{8}{9} \approx 0.888889$
	$+\sqrt{\frac{3}{5}} \approx 0.774597$	$\frac{5}{9} \approx 0.555556$
4	-0.861136	0.347855
	-0.339981	0.652145
	$+0.339981$	0.652145
	$+0.861136$	0.347855

Application to Arc Length. In our arc-length integral, set

$$f(\tau) = \|\mathbf{r}'(\tau)\|, \quad a = 0, \quad b = t,$$

and use the above x_i, w_i with

$$\tau_i = \frac{b-a}{2} x_i + \frac{a+b}{2} = \frac{t}{2} x_i + \frac{t}{2},$$

to compute

$$s(t) \approx \frac{t}{2} \sum_{i=1}^n w_i \|\mathbf{r}'(\tau_i)\|.$$

With just $n = 4$ or 5 , this gives a highly accurate approximation of $s(t)$ without needing extremely fine subdivisions.

7 Finding the Next Parameter: Newton–Raphson

To move a small distance Δs along the path, we need to find the new parameter t_{next} such that

$$s(t_{\text{next}}) = s_{\text{current}} + \Delta s.$$

We use Newton–Raphson (a root-finding method) to solve this equation:

$$t_{k+1} = t_k - \frac{s(t_k) - (s_{\text{current}} + \Delta s)}{v(t_k)}.$$

Why Newton’s method? We want to solve

$$F(t) = s(t) - (s_{\text{current}} + \Delta s) = 0$$

for t . Newton–Raphson is an efficient iterative technique that uses the derivative $F'(t) = s'(t) = v(t)$. Starting from an initial guess t_k , we linearize and obtain

$$t_{k+1} = t_k - \frac{F(t_k)}{F'(t_k)} = t_k - \frac{s(t_k) - (s_{\text{current}} + \Delta s)}{v(t_k)}.$$

Stopping criterion. We iterate $k = 0, 1, 2, \dots$ until one of the following is met:

$$|t_{k+1} - t_k| < \varepsilon \quad \text{or} \quad |s(t_k) - (s_{\text{current}} + \Delta s)| < \varepsilon \quad \text{or} \quad k \geq k_{\text{max}},$$

where $\varepsilon > 0$ is a chosen tolerance (e.g. 10^{-6}) and k_{max} (e.g. 10) caps the number of iterations to avoid infinite loops. This ensures we stop once the update or the residual is sufficiently small, or when we’ve reached a safe iteration limit.

8 Acceleration and Deceleration Distances

From basic kinematics:

$$v^2 = v_0^2 + 2a \Delta s \quad \implies \quad \Delta s = \frac{v^2 - v_0^2}{2a},$$

which tells us the distance needed to change speed under constant acceleration a .

- If $a > 0$: Δs is the distance needed to **accelerate** from v_0 up to v .
- If $a < 0$: Δs is the distance needed to **decelerate** (brake) from v_0 down to v .

Why it works. This formula is derived by integrating acceleration over distance:

$$a = \frac{dv}{dt} \quad \Rightarrow \quad v dv = a ds \quad \Rightarrow \quad \int_{v_0}^v v dv = \int_0^{\Delta s} a ds \quad \Rightarrow \quad \frac{v^2 - v_0^2}{2} = a \Delta s,$$

so rearranging gives $\Delta s = (v^2 - v_0^2)/(2a)$.

Example. Accelerating from 0 to 2 m/s with $a = 1 \text{ m/s}^2$:

$$\Delta s = \frac{2^2 - 0^2}{2 \cdot 1} = \frac{4}{2} = 2 \text{ m}.$$

Choosing a values. In practice, the maximum acceleration a_{\max} and braking deceleration (often taken equal in magnitude for simplicity) are determined by the robot's motors and traction. You can estimate these values from motor specifications or experiments. For example, if a robot can accelerate from 0 to 2 m/s in about 2 seconds, then a rough estimate for a_{\max} is 1 m/s². In reality, a_{\max} might decrease as the robot speeds up (due to motor torque limits and drag), and a_{\max} decel might be limited by wheel grip (to prevent skidding).

9 Curvature and Turning Speed

For a 2D path $\mathbf{r}(t) = (x(t), y(t))$, curvature is defined as:

$$\kappa(t) = \frac{|x'(t)y''(t) - y'(t)x''(t)|}{(x'(t)^2 + y'(t)^2)^{3/2}}, \quad R(t) = \frac{1}{\kappa(t)},$$

where $R(t)$ is the radius of curvature at that point on the path.

To ensure safe turning (limiting lateral acceleration), we restrict speed using:

$$v_{\text{curve}}(t) = v_{\max} \cdot \frac{R(t)}{R(t) + \frac{w}{2}},$$

where w is the track width (distance between the left and right wheels).

Why it works. The centripetal acceleration is $a_c = \frac{v^2}{R}$. As turns get tighter (smaller R), lateral forces increase. The formula above reduces speed as needed to avoid skidding or instability:

- When $R \gg w/2$: $v_{\text{curve}} \approx v_{\max}$ (on gentle curves, you can go full speed).
- When $R \rightarrow 0$: $v_{\text{curve}} \rightarrow 0$ (for very sharp turns, you must crawl).

10 Keyframe Velocity Interpolation

Suppose we have keyframe pairs (s_i, v_i) , which define desired velocities v_i at specific arc lengths (distances) s_i along the path. To determine the target velocity at any intermediate position s (where $s_i \leq s \leq s_{i+1}$), we linearly interpolate between the surrounding keyframes:

$$\lambda = \frac{s - s_i}{s_{i+1} - s_i}, \quad v_{\text{interp}}(s) = v_i + \lambda(v_{i+1} - v_i).$$

Why it works. This simple linear interpolation ensures:

- Smooth, gradual changes between keyframe speeds (no sudden jumps).
- A predictable and easily adjustable speed profile.
- Flexibility in motion design (e.g., you can insert slow-down or speed-up points as needed).

11 Velocity Planning Algorithm

At each control cycle (each small time step), the system computes a safe and desired linear velocity v_{desired} as follows:

1. Compute v_{max} : the robot's absolute maximum speed (a constant limit).
2. Compute v_{accel} : the speed limit imposed by acceleration/braking distance (how fast we can go without overshooting the remaining distance). This comes from the kinematic distance formula above.
3. Compute v_{curve} : the speed limit due to curvature at the robot's current position (using the turning formula).
4. Compute v_{interp} : the speed dictated by keyframe interpolation at the current position.

Then take the minimum of all these constraints:

$$v_{\text{desired}} = \min\{v_{\text{max}}, v_{\text{accel}}, v_{\text{curve}}, v_{\text{interp}}\},$$

which will be the commanded speed for this step.

12 Handling Multiple Path Segments

Many paths consist of several curve segments (splines) joined end-to-end. We can apply the motion profiling approach to each segment in sequence, while ensuring smooth transitions between segments.

At the end of one segment, the next segment begins with the robot continuing at whatever speed it reached. In practice, the end of one spline and the start of the next share a common point and velocity (if planned correctly). To handle the transition smoothly:

- If the remaining distance in the current segment is greater than the distance Δs the robot will travel in the next time step, then the robot simply continues within the same segment (it hasn't reached the end yet).
- If Δs would carry the robot past the end of the current segment, we first move it to the end of that segment (using the portion of Δs needed to cover the last bit of that spline). Then, we use the leftover distance to continue into the next segment (starting the next segment's arc length from 0 at its beginning).

This way, the robot seamlessly continues from one spline to the next. The parameter t resets to 0 at the start of the new segment, and we proceed with the same calculations on the new segment. (If using keyframes, you can also set a keyframe at the segment boundary to enforce a specific speed at that point.)

13 Simplified C++ Example

```
// Compute (v, omega) along a Bezier curve
std::pair<float, float> computeProfile(
    const std::vector<Point>& pts, // 4 control points defining the Bezier
    curve
```



```

float maxVel,           // maximum allowed velocity (m/s)
float maxAcc,           // maximum allowed acceleration (m/s^2)
float decelDist,        // distance from end at which to begin
    braking (m)
float initVel,          // initial velocity at the start of the
    profile (m/s)
float exitVel,          // desired velocity at the end of the
    profile (m/s)
const std::vector<Keyframe>& kfPts, // sequence of (s, v) keyframe
    pairs
float dt                // time step for each profile update (s
    )
) {
    // Static variables preserve values between calls
    // t_prev: previous parameter t along the curve [0,1]
    // v_prev: previous output velocity (initialized to initVel)
    static float t_prev = 0;
    static float v_prev = initVel;

    // 1) Compute the arc-length s at the current parameter t_prev
    //     arcLength(...) integrates the curve length from t=0 to t_prev
    float s_curr = arcLength(pts, t_prev);

    // 2) Determine velocity limit due to acceleration/deceleration
    //     remaining: distance left from s_curr to end of curve (
    //         total_length - s_curr)
    float remaining = arcLength(pts, 1.0f) - s_curr;
    //     regionDist: distance we can accelerate (if before decelDist) or
    //         must decelerate (if within decelDist of end)
    float regionDist = (s_curr < decelDist ? s_curr : remaining);
    //     v_accel: maximum speed reachable under accel limits over
    //         regionDist
    float v_accel = std::sqrt(v_prev*v_prev + 2*maxAcc*regionDist);

    // 3) Determine velocity limit from keyframe constraints
    //     find the segment [s0, s1] in kfPts that contains current s_curr
    int idx = findKeyframeIndex(kfPts, s_curr);
    float s0 = kfPts[idx].s, s1 = kfPts[idx+1].s;
    float v0 = kfPts[idx].v, v1 = kfPts[idx+1].v;
    //     interpolate between v0 and v1 based on where s_curr lies
    float alpha = (s_curr - s0) / (s1 - s0);
    float v_kf = v0 + alpha * (v1 - v0);

    // 4) Determine velocity limit from curvature (turning) constraints
    //     kappa: curvature at current t_prev (1/R)
    float kappa = curvature(pts, t_prev);
    float R = 1.0f / kappa; // radius of curvature
    float w = 0.29f; // track width (example value)
    //     v_curve: max speed so lateral acceleration stays within limits
    float v_curve = maxVel * R / (R + w/2);

    // 5) Combine all velocity limits and clamp to maxVel
    //     v_des is the desired velocity we can safely achieve
    float v_des = std::min({ maxVel, v_accel, v_kf, v_curve });

```

```

// 6) Advance along the curve by ds = v_des * dt
float ds      = v_des * dt;
//    findTForS: solves for new t_next such that arcLength(pts, t_next
//    ) - s_curr = ds
float t_next = findTForS(pts, t_prev, ds);

// 7) Compute angular velocity omega
//    curvature at the new point times linear velocity gives turning
//    rate
float omega = curvature(pts, t_next) * v_des;

// 8) Update stored state for next iteration
t_prev = t_next;
v_prev = v_des;

// Return the computed linear and angular velocities
return { v_des, omega };
}

```

How to use these outputs. Send v_{desired} and ω to your drivetrain’s velocity controller (commonly a PID or feedforward velocity controller, which is beyond the scope of this guide). A typical differential-drive conversion is:

$$v_{\text{left}} = v_{\text{desired}} - \frac{\omega w}{2}, \quad v_{\text{right}} = v_{\text{desired}} + \frac{\omega w}{2},$$

where w is the track width. For even better tracking, you can wrap this in a RAMSETE or pure-pursuit controller that uses the robot’s pose feedback to correct errors and ensure convergence to the planned path.

14 Why Acceleration/Deceleration Distance is an Approximation

The constant- a distance formulas we derived assume the robot accelerates and then decelerates in a perfect “trapezoidal” velocity profile (speeding up, cruising at v_{max} , then slowing down). This assumes the robot has a long enough straight run to reach and maintain that top speed.

However, along a curvy path, we impose extra speed limits based on curvature. This often prevents the robot from ever reaching the planned v_{max} on that segment of the path. In other words, the velocity profile gets “cut off” or dips in the middle due to a turn. The robot might have to start slowing down for the turn before it ever fully accelerates to v_{max} . The figure below illustrates this: in the curved case, the robot cannot maintain a flat-top speed profile because it must slow down for the turn.

Because of this effect, the simple distance formulas for accelerating and braking become approximations. In practice, the robot might start decelerating earlier (or not accelerate as much) when a turn is coming up, meaning our calculated distances might slightly overestimate how far it needs to speed up or slow down.

To improve accuracy, we can account for curvature in our motion profile calculations. For example:

- Instead of relying solely on the constant- a formula, we could integrate the actual velocity curve (which already factors in curvature limits) to find the distance covered during acceleration and deceleration.

- We can dynamically adjust target speeds. If a sharp turn is imminent, the algorithm can cap the speed below v_{\max} ahead of time, so that the robot never accelerates beyond what it can comfortably slow down from when it reaches the curve.

15 Ideas for Future Improvement

- **Motor Model:** Acceleration decreases with speed (motors produce less torque at high RPM). For example:

$$a(v) = a_{\max} \left(1 - \frac{v}{v_{\text{free}}} \right),$$

where v_{free} is the motor's free (no-load) speed.

- **Friction and Drag:** Include frictional forces (wheel friction, air resistance) in the model. These forces oppose motion and reduce net acceleration, especially at higher speeds.
- **Full Dynamics:** Model the robot's full dynamics (mass, rotational inertia) for more realistic acceleration behavior.
- **Numerical Refinements:** Use more advanced numerical methods (adaptive integration steps, etc.) to improve the accuracy of the motion profile calculations.

16 Further Learning Resources

- **YouTube - "The Continuity of Splines":** Deep dive into spline curves by Freya Holmér (<https://www.youtube.com/watch?v=jvPPXbo87ds>)
- **Khan Academy:** Derivatives and Integrals (introductory calculus lessons)
- **Feynman Lectures on Physics**, Vol. 1 (insightful treatment of motion in Chapter 8)
- **MIT OCW:** 18.01 Single Variable Calculus (free online course materials)

17 Conclusion

We have shown how to build a motion profiler that:

- Defines a path with cubic Béziars
- Uses calculus and numerical methods to track arc-length
- Computes acceleration, curvature, and keyframe limits
- Steps along the path with Newton–Raphson root-finding
- Outputs (v, ω) for closed-loop control (e.g. feeding into a controller like RAMSETE)

This approach blends math, programming, and robotics in an accessible, step-by-step way. By understanding the underlying methods (and their limitations), a student can extend and refine the system to handle more complex real-world conditions.

References

- [1] Farin, G. (2002). *Curves and Surfaces for CAGD*. Morgan Kaufmann.
- [2] Coulter, R. C. (1992). Implementation of the pure pursuit path tracking algorithm.
- [3] Abramowitz, M., & Stegun, I. A. (1964). *Handbook of Mathematical Functions*.